

VIKRAM DHANWANI

Data Engineer

Pune, India | vickydhanwani11@gmail.com | +91-7350632588 | linkedin.com/in/vikram-dhanwani

PROFESSIONAL SUMMARY

Data Engineer with 5 years of experience architecting and optimizing enterprise-scale cloud data platforms on Azure. Specialized in building high-performance ETL/ELT pipelines using Azure Data Factory, Databricks, and Delta Lake that process 10TB+ monthly while reducing costs and improving reliability. Proven track record of migrating legacy systems to modern lakehouse architectures, implementing CI/CD practices, and delivering data solutions that drive business intelligence for 500+ users. Expert in cost optimization, having delivered \$200K+ in annual savings through strategic cloud migrations and infrastructure improvements.

Microsoft Certified: Azure Administrator Associate | Azure Solutions Architect Expert

TECHNICAL EXPERTISE

Data Engineering	Azure Data Factory, Synapse Analytics, Databricks, Delta Lake, Apache Spark (PySpark), Data Modeling, ETL/ELT Pipelines, Lakehouse Architecture
Cloud & Storage	Azure Data Lake Gen2, Blob Storage, MySQL, Azure SQL Database, Cosmos DB (NoSQL), Azure Analysis Services
Programming	Python, SQL, PowerShell, Scala, DAX
DevOps & Automation	Azure DevOps, CI/CD Pipelines, YAML, Infrastructure as Code (IaC), Azure Functions, Logic Apps
Analytics & BI	Power BI (Premium, Large Datasets), REST APIs, Data Visualization, Semantic Modeling

PROFESSIONAL EXPERIENCE

- BP

Pune, India

Data Engineer

Sep 2025 - Present

Leading high-performance data engineering initiatives focused on optimizing large-scale Spark ingestion and distributed processing within the cloud data ecosystem.

Optimized high-volume Spark ingestion pipelines, slashing data load times by 80% (from 15 hours to 3 hours) by optimizing JDBC read/write options—leveraging parallel partitioning, fetch-size tuning and fine-tuning Databricks cluster settings

Refactored a Scala-based ingestion framework to support dynamic Spark JDBC read/write options via metadata, enabling teams to optimize I/O concurrency and partitioning at runtime without code changes.

Developed a robust validation suite for high-dimensional business data, automating the verification of complex join logic and data flattening processes. Eliminated manual QA bottlenecks by implementing programmatic checks for pivots and cross-entity consistency, ensuring seamless data flow to downstream consumers.
- SLB (Schlumberger)

Pune, India

Data Engineer II

Nov 2023 - Sep 2025

Architected and delivered data engineering solutions for enterprise analytics platform serving 500+ users across multiple business units, driving cost optimization and operational efficiency.

Architected migration from SAP OData to REST API-based ingestion system, achieving near-real-time data freshness (from daily batches) and reducing extraction failures by 85% through improved error handling and retry mechanisms

Re-engineered 90+ legacy Power BI dataflows into Azure Data Factory pipelines with Delta Lake storage, processing 10TB+ monthly while reducing ETL latency by 30% and enabling incremental refresh patterns for 200+ data assets

- ▶ Built comprehensive monitoring and analytics solution for Power BI Premium capacity utilization across 10 production workspaces, providing real-time visibility into resource consumption patterns and enabling **\$12K/month** in infrastructure rightsizing decisions
- ▶ Implemented passwordless authentication framework using Azure Managed Identity for Synapse Analytics and semantic models, eliminating credential rotation overhead and improving security compliance across 200+ data assets
- ▶ Designed lakehouse utilization analytics platform implementing medallion architecture (bronze/silver/gold) that ingests Log Analytics logs and activity events via REST APIs with PowerShell-based paginated extraction, converts to Parquet with temporal partitioning (**80% storage reduction**), and delivers daily executive insights through star schema semantic model
- ▶ Improved semantic model refresh reliability by **95%** (8 failures/day to 2 failures/week) through OAuth2 token refresh automation to prevent expiration during long-running semantic model refreshes (3+ hours), ServiceNow integration via Logic Apps REST API for automated incident creation, reducing incident response time from 2 days to 6 hours

● **Data Engineer**

Mar 2022 - Nov 2023

Led cloud infrastructure optimization initiatives and data platform migrations to improve cost efficiency and performance.

- ▶ Redesigned Azure Analysis Services logging pipeline by migrating ETL logic from dedicated SQL pool to Synapse Analytics serverless compute, reducing SQL pool costs by **25%** and query volume by **35%** through optimized processing patterns
- ▶ Successfully decommissioned 5 Azure Analysis Services instances by migrating 11 tabular models (500M+ rows) to Power BI Premium Large Datasets, eliminating **\$20K monthly** infrastructure costs while maintaining sub-second query performance
- ▶ Optimized business-critical data warehouse queries from 60 minutes to 2 minutes (**97% improvement**) by migrating distribution strategy from round-robin to hash distribution on high-cardinality keys, implementing clustered columnstore indexes, and automating statistics updates in ETL pipeline for sustained performance
- ▶ Developed custom monitoring solution using Azure Functions and Power BI REST APIs to track health metrics across 10 Premium capacities and 8 on-premises gateways, reducing incident detection time by **40%** and enabling proactive capacity management
- ▶ Automated 20 operational workflows using Azure Functions, Logic Apps, and Power Automate, eliminating **10 hours weekly** of manual data refresh, validation, and reporting tasks
- ▶ Deployed Microsoft Purview sensitivity labels across enterprise Power BI tenant with mandatory labeling policies for all published reports, standardizing data classification compliance across organization

● **Associate Data Engineer**

Oct 2020 - Feb 2022

Established DevOps practices and automated deployment pipelines for data infrastructure, driving efficiency and cost optimization.

- ▶ Spearheaded migration from per-user licensing to Enterprise Shared Capacity model for Power BI Premium, consolidating workloads and reducing annual licensing costs by **\$200K** through strategic capacity planning
- ▶ Implemented CI/CD pipelines in Azure DevOps for 50+ data projects (Azure Data Factory, Synapse Analytics, Power BI), reducing deployment time from 2-3 days to 2-4 hours and eliminating manual deployment errors through automated validation
- ▶ Developed centralized infrastructure monitoring dashboard using PowerShell, Azure Data Factory, and Power Platform, providing unified view of pipeline health and reducing troubleshooting time by **15%**

● **Software Development Intern**

May 2019 - Jul 2019

- ▶ Built real-time drilling operations dashboard using Angular and .NET Core, integrating with SQL Server to display live sensor data from field operations and reducing manual data retrieval time by **70%**

EDUCATION

Bachelor of Technology in Information Technology

Aug 2016 - Aug 2020 | CGPA: 8.46

College of Engineering, Pune (COEP)

Capstone Project: Designed and implemented NLP-based text classification system using scikit-learn and Python to automatically categorize police complaints by crime type. Trained on dataset of 5,000+ labeled complaints, achieving 94% accuracy through feature engineering and ensemble methods.

PROFESSIONAL ACHIEVEMENTS

SLB 2024 Q3 Reward of Excellence

Aker JV Dataflow Design

SLB 2023 Q2 Reward of Excellence

Large Datasets & GBS Migration

SLB 2022 H1 Best Team Player Award

Platform Automation Initiatives

SLB 2021 Reward of Excellence 'WoW'

Infrastructure Enhancement & CI/CD

ADDITIONAL INFORMATION

Languages: English (Fluent), Hindi (Native)

Interests: Data Architecture, Distributed Systems, Cloud Cost Optimization